

Grade of Membership Model and Visualization for RNA-seq data using *CountClust*

Kushal K Dey, Chiaowen Joyce Hsiao & Matthew Stephens

Stephens Lab, The University of Chicago

*Corresponding Email: mstephens@uchicago.edu

March 11, 2016

Abstract

Grade of membership or GoM models (also known as admixture models or Latent Dirichlet Allocation”) are a generalization of cluster models that allow each sample to have membership in multiple clusters. It is widely used to model ancestry of individuals in population genetics based on SNP/ microsatellite data and also in natural language processing for modeling documents [1, 3].

This *R* package implements tools to visualize the clusters obtained from fitting topic models using a Structure plot [2] and extract the top features/genes that distinguish the clusters. In presence of known technical or batch effects, the package also allows for correction of these confounding effects.

CountClust version: 0.1.0 ¹

¹This document used the vignette from *Bioconductor* package *DESeq2*, *cellTree* as *knitr* template

Contents

1	Introduction	2
2	CountClust Installation	2
3	Data Preparation	3
3.0.1	Deng et al 2014	3
3.0.2	GTEX V6 Brain	3
4	Fitting the GoM Model	4
5	Structure plot visualization	4
6	Cluster Annotations	8
7	Supplementary analysis	9
8	Session Info	12

1 Introduction

In the context of RNA-seq expression (bulk or singlecell seq) data, the grade of membership model allows each sample (usually a tissue sample or a single cell) to have some proportion of its RNA-seq reads coming from each cluster. For typical bulk RNA-seq experiments this assumption can be argued as follows: each tissue sample is a mixture of different cell types, and so clusters could represent cell types (which are determined by the expression patterns of the genes), and the membership of a sample in each cluster could represent the proportions of each cell type present in that sample.

Many software packages available for document clustering are applicable to modeling RNA-seq data. Here, we use the R package `maptpx` [4] to fit these models, and we add functionality for visualizing the results and annotating clusters by their most distinctive genes to help biological interpretation. We also provide additional functionality to correct for batch effects and also compare the outputs from two different grade of membership model fits to the same set of samples but different in terms of feature description or model assumptions.

2 CountClust Installation

CountClust requires the following CRAN-R packages: [maptpx](#), [slam](#), [ggplot2](#), [cowplot](#), [parallel](#) along with the *Bioconductor* package: [limma](#).

Installing *CountClust* from *Bioconductor* will install all these dependencies:

```
source("http://bioconductor.org/biocLite.R")
biocLite("CountClust")
```

For installing the working version of this package and loading the data required for this vignette, we use CRAN-R package *devtools*.

```
library(devtools)
install_github('kkdey/CountClust')
```

Then load the package with:

```
library(CountClust)
```

3 Data Preparation

We install the data packages as expressionSet objects for bulk-RNA reads data from GTEx (Genotype Tissue Expression) V6 Project Brain tissue samples [7] and a singlecell-RNA reads data due to Deng *et al* 2014 [6].

```
library(devtools)
install_github('kkdey/singleCellRNASeqMouseDeng2014')
install_github('kkdey/GTExV6Brain')
```

3.0.1 Deng *et al* 2014

Load the scRNA-seq data due to Deng *et al* 2014.

```
library(singleCellRNASeqMouseDeng2014)
deng.counts <- exprs(Deng2014MouseESC)
deng.meta_data <- pData(Deng2014MouseESC)
deng.gene_names <- rownames(deng.counts)
```

3.0.2 GTEx V6 Brain

Load the bulk-RNA seq data from GTEx V6 brain data.

```
library(GTExV6Brain)
gtex.counts <- exprs(GTExV6Brain)
gtex.meta_data <- pData(GTExV6Brain)
gtex.gene_names <- rownames(gtex.counts)
```

4 Fitting the GoM Model

We use a wrapper function of the `topics()` function in the *maptpx* due to Matt Taddy [4].

As an example, we fit the topic model for $k=4$ on the GTEx V6 Brain data and save the GoM model output file to user-defined directory.

```
StructureObj(t(gtex.counts),
             nclus_vec=4, tol=0.1,
             path_rda="./data/MouseDeng2014-topicFit.rda")
```

One can also input a vector of clusters under `nclus_vec` as we do for a list of cluster numbers from 2 to 7 for Deng *et al* 2014 data.

```
StructureObj(t(deng.counts),
             nclus_vec=2:7, tol=0.1,
             path_rda="./data/MouseDeng2014-topicFit.rda")
```

5 Structure plot visualization

Now we perform the visualization for $k=6$ for the Deng *et al* 2014 data.

```
MouseDeng2014.topicFit <- get(load("./data/MouseDeng2014-topicFit.rda"))
names(MouseDeng2014.topicFit$clust_6)
## [1] "K"      "theta" "omega" "BF"     "D"      "X"

omega <- MouseDeng2014.topicFit$clust_6$omega

annotation <- data.frame(
  sample_id = paste0("X", c(1:NROW(omega))),
  tissue_label = factor(rownames(omega),
                        levels = rev( c("zy", "early2cell",
                                         "mid2cell", "late2cell",
                                         "4cell", "8cell", "16cell",
                                         "earlyblast", "midblast",
                                         "lateblast") ) ) )

rownames(omega) <- annotation$sample_id;
```

In the above plot, the samples in each batch have been sorted by the proportional membership of the most representative cluster in that batch. One can also use `order_sample=FALSE` for the un-ordered version, which retains the order as in the data (see Supplementary analysis for example).

Now we perform the Structure plot visualization for $k=4$ for GTEx V6 data for Brain samples .


```
omega <- read.table("../data/omega_cis_genes_brain_2.txt",
                    header = TRUE, sep = " ",
                    stringsAsFactors = FALSE)
dim(omega)
## [1] 1259    4
colnames(omega) <- c(1:NCOL(omega))

tissue_labels <- gtex.meta_data[,3];

annotation <- data.frame(
  sample_id = paste0("X", 1:length(tissue_labels)),
  tissue_label = factor(tissue_labels,
                        levels = rev(unique(tissue_labels) ) ) );

cols <- c("blue", "darkgoldenrod1", "cyan", "red")
```


```
StructureGGplot(omega = omega,
  annotation = annotation,
  palette = RColorBrewer::brewer.pal(8, "Accent"),
  yaxis_label = "Amplification batch",
  order_sample = TRUE,
  axis_tick = list(axis_ticks_length = .1,
    axis_ticks_lwd_y = .1,
    axis_ticks_lwd_x = .1,
    axis_label_size = 7,
    axis_label_face = "bold"))
```



```

StructureGGplot(omega = omega,
  annotation= annotation,
  palette = cols,
  yaxis_label = "",
  order_sample = TRUE,
  split_line = list(split_lwd = .4,
    split_col = "white"),
  axis_tick = list(axis_ticks_length = .1,
    axis_ticks_lwd_y = .1,
    axis_ticks_lwd_x = .1,
    axis_label_size = 7,
    axis_label_face = "bold"))

```



6 Cluster Annotations

We extract the top genes driving each cluster using the `ExtractTopFeatures()` functionality of the *CountClust* package. We first perform the cluster annotations from the GoM model fit with $k=6$ on the single cell RNA-seq data due to Deng *et al*

```
theta_mat <- MouseDeng2014.topicFit$clust_6$theta;
top_features <- ExtractTopFeatures(theta_mat, top_features=100,
                                   method="poisson", options="min");
gene_list <- do.call(rbind, lapply(1:dim(top_features)[1],
                                   function(x) deng.gene_names[top_features[x,]]))
```

We tabulate the top 5 genes for these 6 clusters.

```
library(xtable)
xtable(gene_list[,1:5])
```

	1	2	3	4	5
1	Timd2	Isyna1	Alpl2	Pramel5	Hsp90ab1
2	Upp1	Tdgf1	Aqp8	Fabp5	Tat
3	Actb	Krt18	Fabp3	Id2	Tspan8
4	Rtn2	Ebna1bp2	Zfp259	Nasp	Cenpe
5	LOC100502936	Bcl2l10	Tcl1	E330034G19Rik	Oas1d
6	Obox3	Zfp352	Gm8300	Usp17l5	BB287469

We next perform the same for the topic model fit on the GTEx V6 Brain samples data with $k=4$ clusters.

```
theta_mat <- as.matrix(read.table("../data/theta_cis_genes_brain_2.txt"));
top_features <- ExtractTopFeatures(theta_mat, top_features=100,
                                   method="poisson", options="min");
gene_list <- do.call(rbind, lapply(1:dim(top_features)[1],
                                   function(x) gtex.gene_names[top_features[x,]]))
```

The top 3 genes (ensemble IDs) driving these 4 clusters.

```
library(xtable)
xtable(gene_list[,1:3])
```

	1	2	3
1	ENSG00000120885.15	ENSG00000130203.5	ENSG00000131771.9
2	ENSG00000171617.9	ENSG00000160014.12	ENSG00000154146.8
3	ENSG00000112139.10	ENSG00000139899.6	ENSG00000008710.13
4	ENSG00000197971.10	ENSG00000266844.1	ENSG00000237973.1

References

- [1] Pritchard, Jonathan K., Matthew Stephens, and Peter Donnelly. Inference of population structure using multilocus genotype data. *Genetics*. 155.2, 945-959, 200.
- [2] Rosenberg NA, Pritchard JK, Weber JL, Cann HM, Kidd KK, Zhivotovsky LA, Feldman MW. The genetic structure of human populations. *Science*. 298, 2381-2385, 2002.
- [3] Blei DM, Ng AY, Jordan MI. Latent Dirichlet Allocation. *J. Mach. Learn. Res.* 3, 993-1022, 2003.
- [4] Matt Taddy. On Estimation and Selection for Topic Models. *AISTATS 2012, JMLR W&CP* 22.(maptpx R package), 2012.
- [5] Jaitin DA, Kenigsberg E et al. Massively Parallel Single-Cell RNA-Seq for Marker-Free Decomposition of Tissues into Cell Types. *Science*. 343 (6172) 776-779, 2014.
- [6] Deng Q, Ramskold D, Reinius B, Sandberg R. Single-Cell RNA-Seq Reveals Dynamic, Random Monoallelic Gene Expression in Mammalian Cells. *Science*. 343 (6167) 193-196, 2014.
- [7] The GTEx Consortium. The Genotype-Tissue Expression (GTEx) project. *Nature genetics*. 45(6): 580-585. doi:10.1038/ng.2653, 2013.

7 Supplementary analysis

As an additional analysis, we apply the *CountClust* tools on another single-cell RNA-seq data from mouse spleen due to Jaitin *et al* 2014 [5]. The data had technical effects in the form of amplificationbatch which the user may want to correct for.

We first install and load the data.

```
devtools::install_github('jhsiao999/singleCellRNASeqMouseJaitinSpleen')
```

```
library(singleCellRNASeqMouseJaitinSpleen)
jaitin.counts <- exprs(MouseJaitinSpleen)
jaitin.meta_data <- pData(MouseJaitinSpleen)
jaitin.gene_names <- rownames(jaitin.counts)
```

Extracting the non-ERCC genes satisfying some quality measures.

```
ENSG_genes_index <- grep("ERCC", jaitin.gene_names, invert = TRUE)
jaitin.counts_ensg <- jaitin.counts[ENSG_genes_index, ]
filter_genes <- c("M34473", "abParts", "M13680", "Tmsb4x",
                  "S100a4", "B2m", "Atpase6", "Rpl23", "Rps18",
                  "Rpl13", "Rps19", "H2-Ab1", "Rplp1", "Rpl4",
                  "Rps26", "EF437368")
fcounts <- jaitin.counts_ensg[ -match(filter_genes, rownames(jaitin.counts_ensg)), ]
sample_counts <- colSums(fcounts)
```

```
filter_sample_index <- which(jaitin.meta_data$number_of_cells == 1 &
                             jaitin.meta_data$group_name == "CD11c+" &
                             sample_counts > 600)
fcounts.filtered <- fcounts[,filter_sample_index];
```

We filter the metadata likewise

```
jaitin.meta_data_filtered <- jaitin.meta_data[filter_sample_index, ]
```

We fit the GoM model for $k=7$ and plot the Structure plot visualization to show that the amplification batch indeed drives the clustering patterns.

```
StructureObj(t(fcounts),
             nclus_vec=7, tol=0.1,
             path_rda="../data/MouseJaitinSpleen-topicFit.rda")

#data(MouseJaitinSpleen-topicFit)
MouseJaitinSpleen.topicFit <- readRDS("../data/MouseJaitinSpleen-topicFit.rds")
names(MouseJaitinSpleen.topicFit$clust_7)


## [1] "K"      "theta" "omega" "BF"     "D"      "X"

omega <- MouseJaitinSpleen.topicFit$clust_7$omega

amp_batch <- as.numeric(jaitin.meta_data_filtered[, "amplification_batch"])
annotation <- data.frame(
  sample_id = paste0("X", c(1:NROW(omega))),
  tissue_label = factor(amp_batch,
                        levels = rev(sort(unique(amp_batch)))) ) )
```

It seems from the above Structure plot that amplificationbatch drives the clusters. To remove the effect of amplification batch, one can use. For this, we use the BatchCorrectedCounts() functionality of the package.

```
StructureGGplot(omega = omega,
  annotation = annotation,
  palette = RColorBrewer::brewer.pal(9, "Set1"),
  yaxis_label = "Amplification batch",
  order_sample = FALSE,
  axis_tick = list(axis_ticks_length = .1,
    axis_ticks_lwd_y = .1,
    axis_ticks_lwd_x = .1,
    axis_label_size = 7,
    axis_label_face = "bold"))
```



```
batchcorrect.fcounts <- BatchCorrectedCounts(t(fcounts.filtered), amp_batch);  
dim(batchcorrect.fcounts)  
## [1] 1041 20091
```

8 Session Info

```
sessionInfo()  
  
## R version 3.2.2 (2015-08-14)  
## Platform: x86_64-apple-darwin13.4.0 (64-bit)  
## Running under: OS X 10.10.5 (Yosemite)  
##  
## locale:  
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8  
##  
## attached base packages:  
## [1] parallel stats graphics grDevices utils datasets methods  
## [8] base  
##  
## other attached packages:  
## [1] singleCellRNASeqMouseJaitinSpleen_0.99.0  
## [2] cowplot_0.6.0  
## [3] reshape2_1.4.1  
## [4] ggplot2_2.0.0  
## [5] GTEXV6Brain_0.0.1  
## [6] singleCellRNASeqMouseDeng2014_0.99.0  
## [7] Biobase_2.30.0  
## [8] BiocGenerics_0.16.1  
## [9] CountClust_0.1.0  
## [10] knitr_1.12.3  
##  
## loaded via a namespace (and not attached):  
## [1] Rcpp_0.12.3 digest_0.6.9 plyr_1.8.3 grid_3.2.2  
## [5] gtable_0.1.2 formatR_1.2.1 magrittr_1.5 scales_0.3.0  
## [9] evaluate_0.8 highr_0.5.1 stringi_1.0-1 limma_3.26.3  
## [13] labeling_0.3 BiocStyle_1.8.0 RColorBrewer_1.1-2 tools_3.2.2  
## [17] stringr_1.0.0 munsell_0.4.2 colorspace_1.2-6
```